

CS 484/684 Computational Vision

Many thanks for the design of this assignment go to [Towaki Takikawa](https://tovacinni.github.io/) (<https://tovacinni.github.io/>) and [Olga Veksler](https://cs.uwaterloo.ca/~oveksler/) (<https://cs.uwaterloo.ca/~oveksler/>)

Homework Assignment #5 - Supervised Deep Learning for Segmentation

This assignment will test your understanding of applying deep learning by having you apply (fully supervised) deep learning to semantic segmentation, a well studied problem in computer vision. There is one simple theoretical problem 0. The rest is the programming part.

You can get most of the work done using only CPU, however, the use of GPU will be helpful in later parts. Programming and debugging everything up to and including problem 5c should be fine on CPU. You will notice the benefit of GPU mostly in later parts (d-h) of problem 5, but they are mainly implemented and test your code written and debugged earlier. If you do not have a GPU readily accessible to you, we recommend that you use Google Colaboratory to get access to a GPU. Once you are satisfied with your code up to and including 5(c), simply upload this Jupyter Notebook to Google Colaboratory to run the tests in later parts of Problem 5.

Proficiency with PyTorch is required. Working through the PyTorch tutorials will make this assignment significantly easier. <https://pytorch.org/tutorials/> (<https://pytorch.org/tutorials/>)

Problem 0(a)

Consider linear soft-max classifier $\bar{\sigma}(WX)$ for $K = 2$ (see slides 64-65, topic 10), where $X \in R^{m+1}$ is a homogeneous vector representation of m -dimensional feature (or data point). The classifier parameters matrix W consists of two rows representing linear discriminants W_1 and W_2 in R^{m+1} (including the bias). Prove that soft-max classifier $\bar{\sigma}(WX)$ is equivalent to the sigmoid classifier $\sigma((W_2 - W_1)^T X)$ (slide 48, topic 10).

Your Solution:

Type it here, use latex for math formulas.

$$\begin{aligned}\exp(-(W_2 - W_1)^T x) &= \exp(-(W_2^T - W_1^T)x) = \exp((W_1^T - W_2^T)x) \\ \exp((W_1^T - W_2^T)x) > 0 &\implies \frac{1}{1 + \exp(-(W_2 - W_1)^T x)} > \frac{1}{2} \\ 2 &> 1 + \exp(W_1^T x - W_2^T x) \\ 1 &> \exp(W_1^T x - W_2^T x) \\ 0 &> (W_1^T x - W_2^T x) \\ W_2^T x &> W_1^T x\end{aligned}$$

Problem 0(b)

Consider linear classifier $\bar{\sigma}(WX)$ for any given number of classes K , where $X \in R^{m+1}$ is a homogeneous representation of m -dimensional feature vector and W is a matrix of size $K \times (m + 1)$. Specify an equation for a hyperplane in the feature space corresponding to the *decision boundary* between two classes i and j .

HINT1: Decision boundary is the boundary of two design regions in the feature space: one is a set of all features where the classifier prefers class i , i.e. $\bar{\sigma}_i(WX) \geq \bar{\sigma}_j(WX)$, and the other is a set of features where the classifier prefers class j , i.e. $\bar{\sigma}_i(WX) \leq \bar{\sigma}_j(WX)$.

HINT2: any hyperplane in the feature space R^m can be represented by equation $P^T X = 0$ where $P \in R^{m+1}$ is a vector of the hyperplane parameters. Essentially, your solution should specify P based on the parameters of the linear classifier W . The solution should be very simple - this problem is mainly an exercise in terminology that should force you to review the slides.

Your Solution:

Type it here, use latex for math formulas.

$$g_1(x) = W_1^T x$$

$$g_2(x) = W_2^T x$$

$$g_1(x) - g_2(x) = W_1^T x - W_2^T x = 0$$

$$g_1(x) - g_2(x) = (W_1^T - W_2^T)x = 0$$

$$P = (W_1^T - W_2^T)$$

Programming part

In [79]:

```
%matplotlib inline

# It is best to start with USE_GPU = False (implying CPU). Switch USE_
# we strongly recommend to wait until you are absolutely sure your CPU
USE_GPU = False
```

In [80]:

```
# Python Libraries
import random
import math
import numbers
import platform
import copy

# Importing essential libraries for basic image manipulations.
import numpy as np
import PIL
from PIL import Image, ImageOps
import matplotlib.pyplot as plt
```

In [81]:

```
# We import some of the main PyTorch and TorchVision libraries used for
# Detailed installation instructions are here: https://pytorch.org/get-started/locally/
# That web site should help you to select the right 'conda install' command
# In particular, select the right version of CUDA. Note that prior to
# install the latest driver for your GPU and CUDA (9.2 or 10.1), assuming
# For more information about pytorch refer to
# https://pytorch.org/docs/stable/torchvision/transforms.html
# https://pytorch.org/docs/stable/data.html
# and https://pytorch.org/docs/stable/torchvision/transforms.html

import torch
import torch.nn.functional as F
from torch import nn
from torch.utils.data import DataLoader
import torchvision.transforms as transforms
import torchvision.transforms.functional as tF

# We provide our own implementation of torchvision.datasets.voc (containing
# that allows us to easily create single-image datasets

# Note class labels used in Pascal dataset:
# 0: background,
# 1-20: aeroplane, bicycle, bird, boat, bottle, bus, car, cat, chair,
#         person, pottedplant, sheep, sofa, train, TV_monitor
# 21-25: diningtable, motorbike, sofa, televisionmonitor, vodoprovod
```

In [82]:

```
# ChainerCV is a library similar to TorchVision, created and maintained by
# Chainer, the base library, inspired and led to the creation of PyTorch
# Although Chainer and PyTorch are different, there are some nice functions
# that are useful, so we include it as an exercise on learning other libraries
# To install ChainerCV, normally it suffices to run "pip install chainercv"
# For more detailed installation instructions, see https://chainercv.readthedocs.io/en/latest/installation.html
# For other information about ChainerCV library, refer to https://chainercv.readthedocs.io/en/latest/index.html

from chainercv.evaluations import eval_semantic_segmentation
```

```
In [84]: # This colorize_mask class takes in a numpy segmentation mask,  
# and then converts it to a PIL Image for visualization.  
# Since by default the numpy matrix contains integers from  
# 0,1,...,num_classes, we need to apply some color to this  
# so we can visualize easier! Refer to:  
# https://pillow.readthedocs.io/en/4.1.x/reference/Image.html#PIL.Ima  
palette = [0, 0, 0, 128, 0, 0, 0, 128, 0, 128, 128, 0, 0, 0, 128, 128,  
          128, 128, 128, 64, 0, 0, 192, 0, 0, 64, 128, 0, 192, 128, 0,  
          64, 128, 128, 192, 128, 128, 0, 64, 0, 128, 64, 0, 0, 192,  
  
def colorize_mask(mask):  
    new_mask = Image.fromarray(mask.astype(np.uint8)).convert('P')  
    new_mask.putpalette(palette)
```

```
In [85]:
```

```
# Below we will use a sample image-target pair from VOC training dataset
# Running this block will automatically download the PASCAL VOC Dataset
# The code below creates subdirectory "datasets" in the same location as this notebook
# you can modify DATASET_PATH to download the dataset to any custom directory
# On subsequent runs you may save time by setting "download = False" (True is the default)

DATASET_PATH = 'datasets'

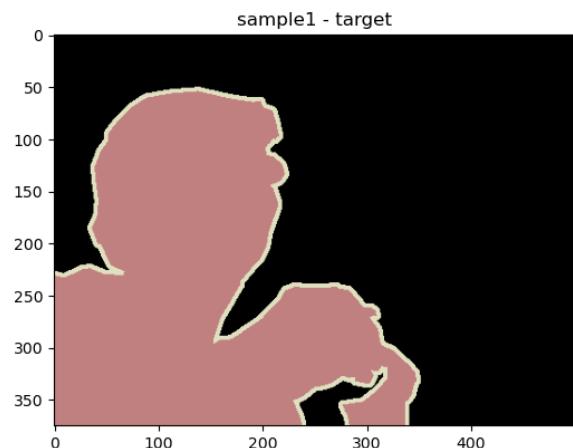
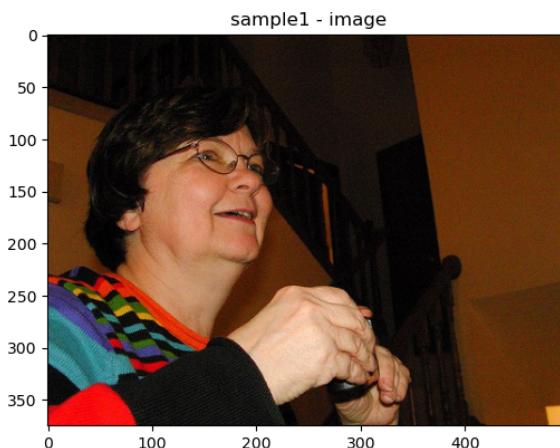
# Here, we obtain and visualize one sample (img, target) pair from VOC
# Note that operator [...] extracts the sample corresponding to the specified index
# Also, note the parameter download = True. Set this to False after you have downloaded the dataset
sample1 = VOCSegmentation(DATASET_PATH, image_set='train', download = True)
sample2 = VOCSegmentation(DATASET_PATH, image_set='val')[20]

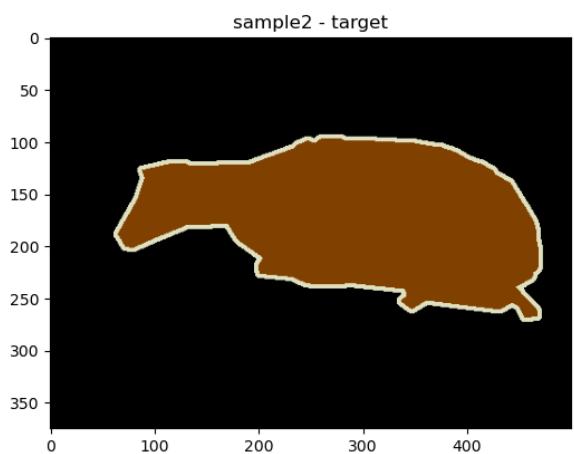
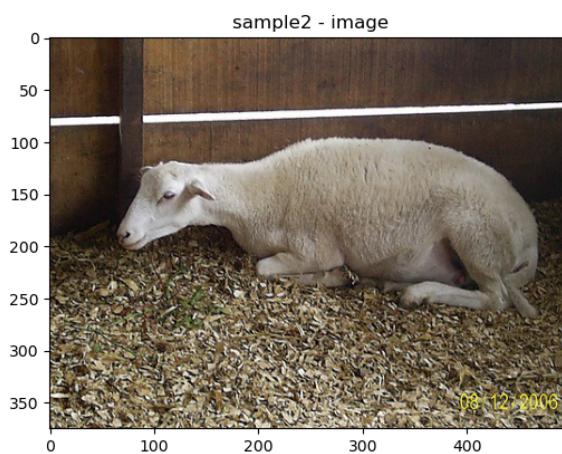
# We demonstrate two different (equivalent) ways to access image and target
img1, target1 = sample1
img2 = sample2[0]
target2 = sample2[1]

fig = plt.figure(figsize=(14,10))
ax1 = fig.add_subplot(2,2,1)
plt.title('sample1 - image')
ax1.imshow(img1)
ax2 = fig.add_subplot(2,2,2)
plt.title('sample1 - target')
ax2.imshow(target1)
ax3 = fig.add_subplot(2,2,3)
plt.title('sample2 - image')
ax3.imshow(img2)
ax4 = fig.add_subplot(2,2,4)
plt.title('sample2 - target')
ax4.imshow(target2)
```

Using downloaded and verified file: datasets/VOCtrainval_11-May-2012.tar

Out [85]: <matplotlib.image.AxesImage at 0x14fc89d0>





Problem 1

Implement a set of "Joint Transform" functions to perform data augmentation in your dataset.

Neural networks are typically applied to transformed images. There are several important reasons for this:

1. The image data should be in a certain required format (i.e. consistent spatial resolution to batch). The images should also be normalized and converted to the "tensor" data format expected by pytorch libraries.
2. Some transforms are used to perform randomized image domain transformations with the purpose of "data augmentation".

In this exercise, you will implement a set of different transform functions to do both of these things. Note that unlike classification nets, training semantic segmentation networks requires that some of the transforms are applied to both image and the corresponding "target" (Ground Truth segmentation mask). We refer to such transforms and their compositions as "Joint". In general, your Transform classes should take as the input both the image and the target, and return a tuple of the transformed input image and target. Be sure to use critical thinking to determine if you can apply the same transform function to both the input and the output.

For this problem you may use any of the `torchvision.transforms.functional` functions. For inspiration, refer to:

https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
[\(https://pytorch.org/tutorials/beginner/data_loading_tutorial.html\)](https://pytorch.org/tutorials/beginner/data_loading_tutorial.html)

<https://pytorch.org/docs/stable/torchvision/transforms.html#module-torchvision.transforms.functional>
[\(https://pytorch.org/docs/stable/torchvision/transforms.html#module-torchvision.transforms.functional\)](https://pytorch.org/docs/stable/torchvision/transforms.html#module-torchvision.transforms.functional)

Example 1

This class takes a img, target pair, and then transform the pair such that they are in `Torch.Tensor()` format.

Solution:

```
In [86]: class JointToTensor(object):
    def __call__(self, img, target):
        return tF.to_tensor(img), torch.from_numpy(np.array(target.con
```

```
In [87]: # Check the transform by passing the image-target sample.
```

```
img, target = JointToTensor()(*sample1)
print(type(img))
print(type(target))
print(img.shape)
print(target.shape)

<class 'torch.Tensor'>
<class 'torch.Tensor'>
torch.Size([3, 375, 500])
torch.Size([375, 500])
```

Example 2:

This class implements CenterCrop that takes an img, target pair, and then apply a crop about the center of the image such that the output resolution is size × size.

Solution:

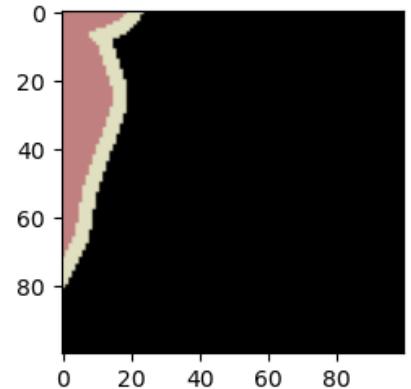
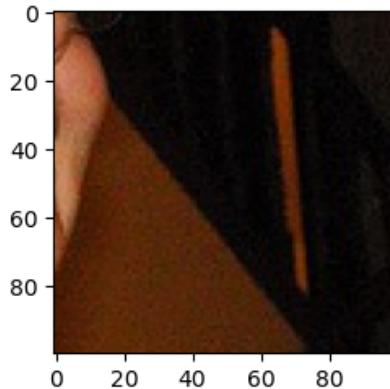
```
In [88]: class JointCenterCrop(object):
    def __init__(self, size):
        """
        params:
            size (int) : size of the center crop
        """
        self.size = size

    def __call__(self, img, target):
        return (tF.five_crop(img, self.size)[4],
                tF.five_crop(target, self.size)[4])

img, target = JointCenterCrop(100)(*sample1)
print(type(img))
print(type(target))
fig = plt.figure(figsize=(12,6))
ax1 = fig.add_subplot(2,2,1)
ax1.imshow(img)
ax2 = fig.add_subplot(2,2,2)
ax2.imshow(target)
```

```
<class 'PIL.Image.Image'>
<class 'PIL.Image.Image'>
```

Out[88]: <matplotlib.image.AxesImage at 0x15cf1af10>



(a) Implement RandomFlip

This class should take a img, target pair and then apply a horizontal flip across the vertical axis at random.

Solution:

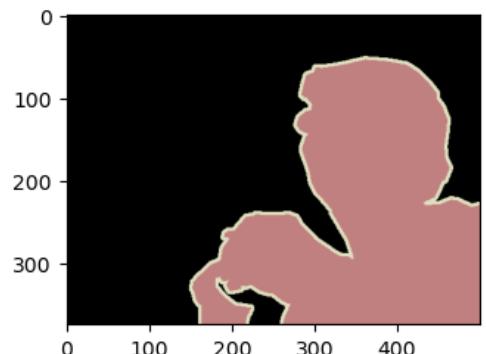
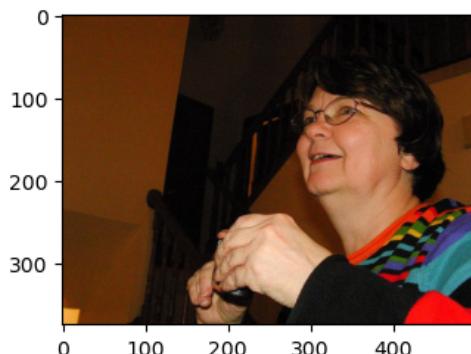
In [89]:

```
class RandomFlip(object):
    def __init__(self):
        pass

    def __call__(self, img, target):
        if torch.rand(1)>0.5:
            return (tF.hflip(img),
                    tF.hflip(target))
        else:
            return (img, target)

img, target = RandomFlip()(*sample1)
fig = plt.figure(figsize=(12,6))
ax1 = fig.add_subplot(2,2,1)
ax1.imshow(img)
ax2 = fig.add_subplot(2,2,2)
ax2.imshow(target)
```

Out[89]: <matplotlib.image.AxesImage at 0x15ce87c70>



(b) Implement RandomResizeCrop

This class should take a img, target pair and then resize the images by a random scale between [minimum_scale, maximum_scale], crop a random location of the image by min(size, image_height, image_width) where the size is passed in as an integer in the constructor, and then resize to size × size (again, the size passed in). The crop box should fit within the image.

Solution:

In [90]: `import random`

This class should take a img, target pair and then resize the images by crop a random location of the image by $\min(\text{size}, \text{image_size})$ in the constructor), and then resize to $\text{size} \times \text{size}$

```
class Resize(object):
    def __init__(self, size_height=100, size_width=100):
        self.size_height = size_height
        self.size_width = size_width
        pass

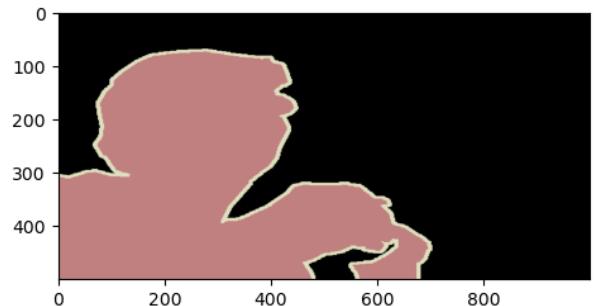
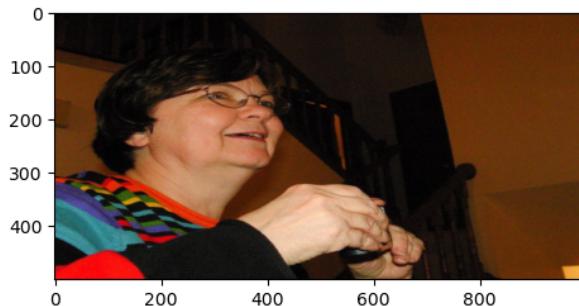
    def __call__(self, img, target):
        size_height = self.size_height
        size_width = self.size_width
        img = tf.resize(img, (size_height, size_width))
        target = tf.resize(target, (size_height, size_width))

        return (img, target)

image_input_size_new = (500, 1000)

img, target = Resize(image_input_size_new[0], image_input_size_new[1])()
fig = plt.figure(figsize=(12, 6))
x1 = fig.add_subplot(2, 2, 1)
x1.imshow(img)
x2 = fig.add_subplot(2, 2, 2)
x2.imshow(target)
```

Out[90]: <matplotlib.image.AxesImage at 0x15cebb820>



```
In [91]: import random
#This class should take a img, target pair and then resize the images &
#crop a random location of the image by $\min(\lceil \frac{size}{image\_size} \rceil, size) times
#in the constructor), and then resize to $\lceil \frac{size}{image\_size} \rceil \times \lceil \frac{size}{image\_size} \rceil * size$.

class RandomResizeCrop(object):
    def __init__(self,min_scale=0.9,max_scale=1.1, size=100):
        self.min_scale = min_scale
        self.max_scale = max_scale
        self.size = size
        pass

    def __call__(self, img, target):
        scale = torch.rand(1)*(self.max_scale-self.min_scale)+self.min_
        w,h=img.size[-2:]
        new_scale = (int(h*scale), int(w*scale))
        resized_img = tF.resize(img, new_scale)
        resized_tar = tF.resize(target, new_scale)
        size = int(min(self.size, new_scale[0], new_scale[1]))
        x_crop = random.randint(0, new_scale[1]-size)
        y_crop = random.randint(0, new_scale[0]-size)
        cropped_img = tF.crop(resized_img, y_crop, x_crop, size,size)
        cropped_tar = tF.crop(resized_tar, y_crop, x_crop, size,size)

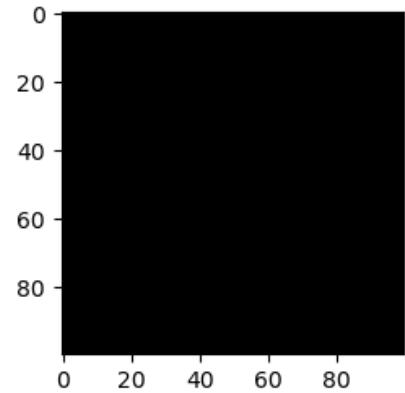
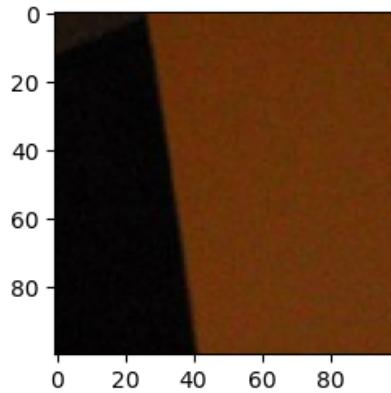
        cropped_img = tF.resize(cropped_img, (size,size))
        cropped_tar = tF.resize(cropped_tar, (size,size))

        return (cropped_img, cropped_tar)
```

```
In [92]: print(sample1[0].size)
      img, target = RandomResizeCrop(min_scale=0.9, max_scale=1.1, size=100)
      fig = plt.figure(figsize=(12,6))
      ax1 = fig.add_subplot(2,2,1)
      ax1.imshow(img)
      ax2 = fig.add_subplot(2,2,2)
      ax2.imshow(target)
```

(500, 375)

Out[92]: <matplotlib.image.AxesImage at 0x15cfb8370>



(c) Implement Normalize

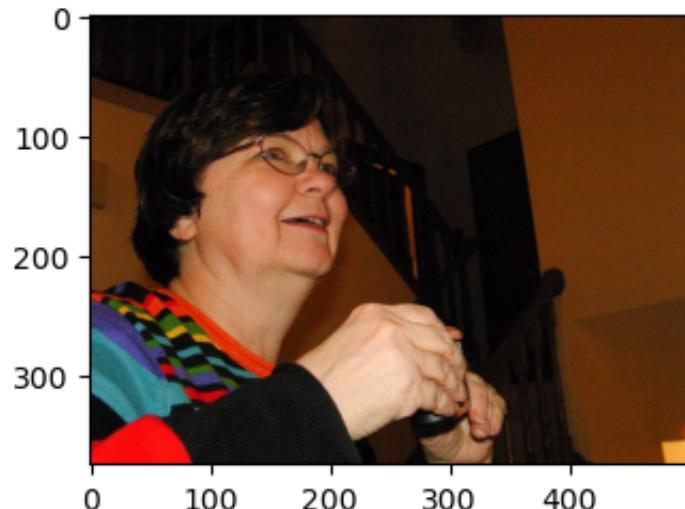
In [93]:

```
class NullTransform(object):
    def __init__(self):
        """
        params:
            p (float) : probability of the horizontal flip
        """

    def __call__(self, img, target):
        return img, target

img, target = NullTransform()(*sample1)
fig = plt.figure(figsize=(12,6))
ax1 = fig.add_subplot(2,2,1)
ax1.imshow(img)
```

Out[93]: <matplotlib.image.AxesImage at 0x15d0325e0>



In [94]: type(img)

Out[94]: PIL.Image.Image

In [95]:

```
norm = ([0.485, 0.456, 0.406],
         [0.229, 0.224, 0.225])

class Normalize(object):
    def __init__(self, mean, std):
        self.mean = mean
        self.std = std

    def __call__(self, tensor, target):
        return (tF.normalize(tensor, self.mean, self.std), target)

tensor = JointToTensor()(*sample1)
Normalize(*norm)(*tensor)
```

Out[95]:

```
(tensor([[[ -1.9295, -1.8953, -1.9638, ..., -0.7479, -0.4911, -0.5938
        ],
                [-2.0323, -1.9295, -2.0152, ..., -0.6794, -0.5596, -0.6794
        ],
                [-1.9467, -1.8953, -1.9124, ..., -0.6281, -0.4739, -0.5424
        ],
                ...,
                [ 1.9920,  2.2318,  2.2489, ...,  2.0948,  2.1119,  2.1290
        ],
                [ 2.2489,  2.1119,  2.2318, ...,  2.0777,  2.0948,  2.1119
        ],
                [ 2.2489,  2.0263,  2.2489, ...,  2.1290,  2.1290,  2.1804
        ]],
        [[[ -1.8606, -1.8256, -1.9482, ..., -1.5105, -1.2479, -1.3004
        ],
                [-1.9657, -1.8606, -2.0007, ..., -1.4580, -1.3179, -1.4405
        ],
                [-1.8782, -1.8256, -1.8957, ..., -1.4580, -1.2479, -1.3179
        ],
                ...,
                [-1.9657, -1.7381, -1.7206, ...,  0.8880,  0.9055,  0.9055
        ],
                [-1.6856, -1.8431, -1.7381, ...,  0.8704,  0.8704,  0.8880
        ],
                [-1.7031, -1.9307, -1.6856, ...,  0.9230,  0.9230,  0.9580
        ]],
        [[[ -1.6650, -1.6302, -1.7347, ..., -1.7522, -1.4907, -1.5604
        ],
                [-1.7696, -1.6650, -1.7870, ..., -1.6999, -1.5604, -1.6824
        ],
                [-1.6824, -1.6302, -1.6824, ..., -1.6824, -1.4907, -1.5604
        ]],
```

```
        ...,
        [-1.5604, -1.3687, -1.3861, ..., -0.5495, -0.5321, -0.5321
],
        [-1.2816, -1.4384, -1.3687, ..., -0.5670, -0.5670, -0.5495
],
        [-1.2641, -1.5256, -1.3164, ..., -0.5147, -0.5147, -0.4798
]]]),  
    tensor([[ 0,  0,  0, ...,  0,  0,  0],  
           [ 0,  0,  0, ...,  0,  0,  0],  
           [ 0,  0,  0, ...,  0,  0,  0],  
           ...,  
           [15, 15, 15, ...,  0,  0,  0],  
           [15, 15, 15, ...,  0,  0,  0],  
           [15, 15, 15, ...,  0,  0,  0]]))
```

This class should take a img, target pair and then normalize the images by subtracting the mean and dividing variance.

Solution:**(d) Compose the transforms together:**

Use `JointCompose` (fully implemented below) to compose the implemented transforms together in some random order. Verify the output makes sense and visualize it.

In [96]: This class composes transformations from a given list of image transformations that will be applied to the dataset during training. This cell is fully implemented.

```
class JointCompose(object):
    def __init__(self, transforms):
        """
        params:
            transforms (list) : list of transforms
        """
        self.transforms = transforms

    # We override the __call__ function such that this class can be
    # called as a function i.e. JointCompose(transforms)(img, target)
    # Such classes are known as "functors"
    def __call__(self, img, target):
        """
        params:
            img (PIL.Image)      : input image
            target (PIL.Image)   : ground truth label
        """
        assert img.size == target.size
        for t in self.transforms:
            img, target = t(img, target)
        return img, target
```

In [97]: Student Answer:

This class composes transformations from a given list of image transformations that will be applied to the dataset during training. This cell is fully implemented.

```
class JointComposeRandom(object):
    def __init__(self, transforms):
        """
        params:
            transforms (list) : list of transforms
        """
        self.transforms = transforms

    # We override the __call__ function such that this class can be
    # called as a function i.e. JointCompose(transforms)(img, target)
    # Such classes are known as "functors"
    def __call__(self, img, target):
        random.shuffle(self.transforms)
        JointCompose(self.transforms)(img, target)
        return img, target
```

(e) Compose the transforms together: use `JointCompose` to compose the implemented transforms for:

- 1. A sanity dataset that will contain 1 single image. Your objective is to overfit on this 1 image, so choose your transforms and parameters accordingly.**
- 2. A training dataset that will contain the training images. The goal here is to generalize to the validation set, which is unseen.**
- 3. A validation dataset that will contain the validation images. The goal here is to measure the 'true' performance.**

In [98]:

```
norm = ([0.485, 0.456, 0.406],  
        [0.229, 0.224, 0.225])
```

In [99]:

```
Image input size: (275, 500)
```

In [100]: # Student Answer:

```
sanity_joint_transform = JointCompose([ \
    RandomResizeCrop(0.9, 2, 250),
    JointCenterCrop(224),
    RandomFlip(),
    Resize(image_input_size[0], image_input_size[1]),
    JointToTensor(),
    Normalize(*norm),
])

train_joint_transform = JointCompose([ \
    Resize(image_input_size[0], image_input_size[1]),
    RandomResizeCrop(0.9, 2, 250),
    JointCenterCrop(224),
    RandomFlip(),
    Resize(image_input_size[0], image_input_size[1]),
    JointToTensor(),
    Normalize(*norm),
])

val_joint_transform = JointCompose([ \
    Resize(image_input_size_new[0], image_input_size_new[1]),
    RandomResizeCrop(1, 1, 250),
    JointCenterCrop(224),
    RandomFlip(),
    Resize(image_input_size[0], image_input_size[1]),
    JointToTensor(),
    Normalize(*norm),
])
```

Student Answer:

```
sanity_joint_transform = JointCompose([\
```

RandomResizeCrop(0.9,2, 250),

JointCenterCrop(224),

RandomFlip(),

```
    JointToTensor(),
    Normalize(*norm),
```

```
])
```

```
train_joint_transform = JointCompose([
    RandomResizeCrop(0.9,2, 250), JointCenterCrop(224), RandomFlip(), JointToTensor(),
    Normalize(*norm), ])
```

```
val_joint_transform = JointCompose([
    RandomResizeCrop(1,1, 250),
```

JointCenterCrop(224),

RandomFlip(),

```
    JointToTensor(),
    Normalize(*norm),
```

```
])
```

This code below will then apply `train_joint_transform` to the entire dataset.

In [101]: *Copy the Joint-Compose transformations above to create three datasets as this cell is fully implemented.*

This single image data(sub)set can help to better understand and to debug. The optional integer parameter 'sanity_check' specifies the index of the image to be checked. Note that we use the same image (index=200) as used for sample1.

```
sanity_data = VOCSegmentation(  
    DATASET_PATH,  
    image_set = 'train',  
    transforms = sanity_joint_transform,  
    sanity_check = 200)
```

This is a standard VOC data(sub)set used for training semantic segmentation.

```
train_data = VOCSegmentation(  
    DATASET_PATH,  
    image_set = 'train',  
    transforms = train_joint_transform)
```

This is a standard VOC data(sub)set used for validating semantic segmentation.

```
val_data = VOCSegmentation(  
    DATASET_PATH,  
    image_set='val',  
    transforms = val_joint_transform)
```

Increase TRAIN_BATCH_SIZE if you are using GPU to speed up training. When batch size changes, the learning rate may also need to be adjusted. Note that batch size maybe limited by your GPU memory, so adjust if you like.

```
TRAIN_BATCH_SIZE = 4
```

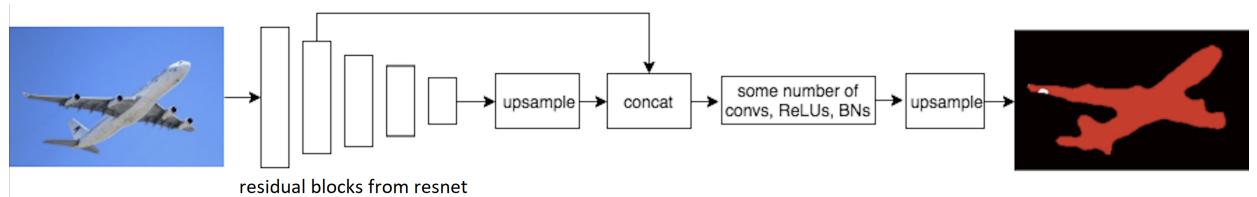
If you are NOT using Windows, set NUM_WORKERS to anything you want, e.g. 4. But Windows has issues with multi-process dataloaders, so NUM_WORKERS must be 0.

```
sanity_loader = DataLoader(sanity_data, batch_size=1, num_workers=NUM_WORKERS)  
train_loader = DataLoader(train_data, batch_size=TRAIN_BATCH_SIZE, num_workers=NUM_WORKERS)  
val_loader = DataLoader(val_data, batch_size=1, num_workers=NUM_WORKERS, shuffle=False)
```

Problem 2

(a) Implement encoder/decoder segmentation CNN using PyTorch.

You must follow the general network architecture specified in the image below. Note that since convolutional layers are the main building blocks in common network architectures for image analysis, the corresponding blocks are typically unlabeled in the network diagrams. The network should have 5 (pre-trained) convolutional layers (residual blocks) from "resnet" in the encoder part, two upsampling layers, and one skip connection. For the layer before the final upsampling layer, lightly experiment with some combination of Conv, ReLU, BatchNorm, and/or other layers to see how it affects performance.



You should choose specific parameters for all layers, but the overall structure should be restricted to what is shown in the illustration above. For inspiration, you can refer to papers in the citation section of the following link to DeepLab (e.g. specific parameters for each layer): [\(http://liangchihchen.com/projects/DeepLab.html\)](http://liangchihchen.com/projects/DeepLab.html). The first two papers in the citation section are particularly relevant.

In your implementation, you can use a base model of choice (you can use `torchvision.models` as a starting point), but we suggest that you learn the properties of each base model and choose one according to the computational resources available to you.

Note: do not apply any post-processing (such as DenseCRF) to the output of your net.

Solution:

```
import torchvision.models as models

class MyNet(nn.Module): def __init__(self, num_classes, criterion=None): super(MyNet, self).__init__()

    self.net = nn.Sequential(
        nn.Linear(2, 4),
        nn.ReLU(),
        nn.Linear(4, 4),
        nn.ReLU(),
        nn.Linear(4, 2),
        nn.Softmax(dim=1),
    )

    def forward(self, inp, gts=None):

        # Implement me

        if self.training:
            # Return the loss if in training mode
            return self.criterion(lfinal, gts)
        else:
            # Return the actual prediction otherwise
            return lfinal
```

In [102]: `import torchvision.models as models`

```
class MyNet(nn.Module):
    def __init__(self, num_classes, criterion=None):
        super(MyNet, self).__init__()
        self.resnet_fst_layer = nn.Sequential(*[list(models.resnet18(precise_bn=True).children())[:5]])
        self.resnet_other_layers = nn.Sequential(*[list(models.resnet18(precise_bn=True).children())[5:]])
        self.conv1 = nn.ConvTranspose2d(64+512, 400, 5, 1)
        self.conv2 = nn.ConvTranspose2d(400, 512, 3, 1)
        self.conv3 = nn.ConvTranspose2d(512, 256, 3, 1)
        self.conv4 = nn.Conv2d(256, num_classes, 1, 1)
        self.BN1 = nn.BatchNorm2d(256)
        self.BN2 = nn.BatchNorm2d(512)

    def forward(self, inp, gts=None):
        x = inp
        # print(x.size())
        l1 = x = self.resnet_fst_layer(x)
        # print(x.size())
        x = self.resnet_other_layers(x)
        x = F.interpolate(x, l1.size()[-2:])
        x = torch.cat((x, l1), dim=1)

        x = F.relu(self.conv1(x))
        x = F.relu(self.BN2(self.conv2(x)))
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = F.interpolate(x, inp.size()[-2:])
        lfinal = x

        if self.training:
            # Return the loss if in training mode
            return self.criterion(lfinal, gts)
        else:
            # Return the actual prediction otherwise
            return lfinal
```

(b) Create UNTRAINED_NET and run on a sample image

```
In [103]: untrained_net = MyNet(21).eval()
sample_img, sample_target = Normalize(*norm)(*JointToTensor())(*sample1
untrained_output = untrained_net.forward(sample_img[None])
```

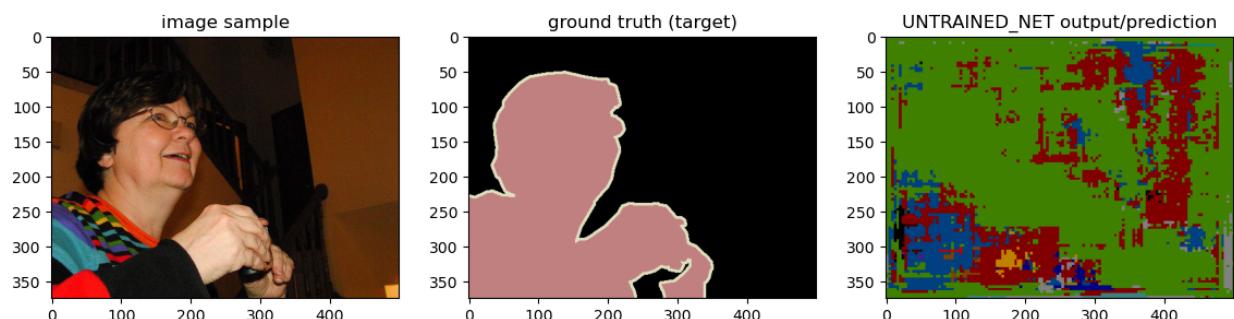
```
In [104]: sample_img.shape
```

```
Out[104]: torch.Size([3, 375, 500])
```

```
In [105]:
```

```
fig = plt.figure(figsize=(14,10))
ax = fig.add_subplot(1,3,1)
plt.title('image sample')
ax.imshow(sample1[0])
ax = fig.add_subplot(1,3,2)
plt.title('ground truth (target)')
ax.imshow(sample1[1])
ax = fig.add_subplot(1,3,3)
plt.title('UNTRAINED_NET output/prediction')
ax.imshow(colorize_mask(torch.argmax(untrained_output, dim=1).numpy())]
```

```
Out[105]: <matplotlib.image.AxesImage at 0x15d18c490>
```



Problem 3 (bonus for undergrads, required for grads)

(a) Implement the loss function (Cross Entropy Loss) as detailed below. For debugging, part (b) below allows to compare the values w.r.t. the standard function "nn.CrossEntropyLoss". If your loss function works, uncomment the use of "MyCrossEntropyLossy" replacing "nn.CrossEntropyLoss" everywhere below in the notebook (just a couple of places).

You should return the mean (average over all batch pixels) negative log of soft-max for the ground truth class. Assuming that

X_p^m are logits at pixel p for C classes, so that $m \in \{0, 1, \dots, C - 1\}$, each point p contributes to this loss

$$-\log \frac{e^{X_p^{y_p}}}{\sum_m e^{X_p^m}} = -\sum_k Y_p^k \log \frac{e^{X_p^k}}{\sum_m e^{X_p^m}} \quad (*)$$

where y_p is the ground truth label and Y_p^k is the corresponding one-hot distribution. You should implement the basic math, as in one of the equations above.

NOTE 1: Numerically robust implementation of soft-max is not immediately straightforward. Indeed, logits X_p^k are arbitrary numbers and their exponents could be astronomically large. While the log cancels out the exponent in the numerator (do it mathematically, not numerically), the sum of exponents in the denominator can not be easily simplified and this is where the numerical issues hide. A naive implementation adding huge values of logits' exponents easily loses precision. Instead, one should use the right hand side in the following algebraically equivalent formulation:

$$\log(e^A + e^B + \dots + e^Z) \equiv \mu + \log(e^{A-\mu} + e^{B-\mu} + \dots + e^{Z-\mu})$$

where $\mu := \max\{A, B, \dots, Z\}$. The right hand side is numerically stable since the arguments of the exponents are negative and the exponents are all bounded by value 1.

HINT 1: Similarly to many previous assignments, avoid for-loops. In fact, you should not use none below. Instead, use standard operators or functions for tensors/matrices (e.g. pointwise addition, multiplication, etc.) In particular, you will find useful functions like "sum" and "max" that could be applied along any specified dimension of the tensor. Of course, torch also has pointwise functions "log" and "exp" that could be applied to any tensor.

HINT 2: Be careful - you will have to work with tensors of different shapes, e.g. even input tensors "targets" and "logits" have different shapes. As in previous assignments based on "numpy", you should pay attention to tensors' dimensions in pyTorch. For example, pointwise addition or multiplication requires either the same shape or shapes "broadcastable" to the same shape (similar in "pyTorch" and "numpy", e.g. see [here](#) (<https://numpy.org/doc/stable/user/theory.broadcasting.html#array-broadcasting-in-numpy>)).

If in doubt, use

print(xxx.size())

to check the shape - I always do! If needed, modify the shapes using standard functions: "transpose" to change the order of the dimensions, "reshape", "flatten", or "squeeze" to remove dimensions, "unsqueeze" to add dimensions. You can use other standard ways to modify the shape, or rely on broadcasting.

HINT 3: Your loss should be averaged only over pixels that have non-void labels. That is, exclude pixels with "ignore index" labels. For example, you can compute a "mask" of non-void pixels, and use it to trim non-void pixels in both the "targets" and "logits". You might get an inspiration from the "mask" in your K-means implementation.

HINT 4: For simplicity, you may "flatten" all tensors' dimensions corresponding to batches and image width & height. The loss computation does not depend on any information in these dimensions and all (non-void) pixels contribute independently based on their logits and ground truth labels (as in the formula above).

HINT 5: In case you want to use the right-hand-side in (*), you can use the function "torch.nn.functional.one_hot" to convert a tensor of "target" class labels (integers) to the tensor of one-hot distributions. Note that this adds one more dimension to the tensor. In case you want to implement the left-hand-side of (*), you can use the function "torch.gather" to select the ground truth class logits $X_p^{y_p}$.

HINT 6: Just as some guidance, a good solution should be around ten lines of "basic" code, or less.

In [106]: # Student Answer:

```
class MyCrossEntropyLoss():
    def __init__(self, ignore_index):
        self.ignore_index = ignore_index

    def __call__(self, untrained_output, sample_target):
        softmax = F.log_softmax(untrained_output, dim = 1)
        ...
```

(b) Compare against the existing CrossEntropyLoss function on your sample output from your neural network.

In [107]: criterion = nn.CrossEntropyLoss(ignore_index=255)

print(criterion(untrained_output, sample_target[None]))

my_criterion = MyCrossEntropyLoss(ignore_index=255)

print(my_criterion(untrained_output, sample_target[None]))

```
tensor(3.0495, grad_fn=<NllLoss2DBackward0>)
tensor(3.0495, grad_fn=<NllLoss2DBackward0>)
```

Problem 4

(a) Use standard function `eval_semantic_segmentation` (already imported from chainerCV) to compute "mean intersection over union" for the output of UNTRAINED_NET on sample1 (`untrained_output`) using the target for sample1. Read documentations for function `eval_semantic_segmentation` to properly set its input parameters.

```
In [108]: # Write code to properly compute 'pred' and 'gts' as arguments for func  
  
pred = [torch.argmax(untrained_output, dim=1)[0].numpy()]  
gts = [torch.where(sample_target == 255, -1, sample_target).numpy()]  
  
conf = eval_semantic_segmentation(pred, gts)  
  
mIoU for the sample image / ground truth pair: 9.375e-05
```

(b) Write the validation loop.

```
In [124]: def validate(val_loader, net):  
  
    iou_arr = []  
    val_loss = 0  
  
    with torch.no_grad():  
        for i, data in enumerate(val_loader):  
  
            inputs, masks = data  
  
            if USE_GPU:  
                inputs = inputs.cuda()  
                masks = masks.cuda()  
                net = net.cuda()  
  
            output = net(inputs)  
            val_loss += criterion(output, masks).sum().item()  
  
            pred = torch.argmax(output, dim=1).numpy()  
            gts = torch.where(masks == 255, -1, masks).numpy()  
  
            # Hint: make sure the range of values of the ground truth  
            conf = eval_semantic_segmentation(pred, gts)  
  
            iou_arr.append(conf['miou'])
```

(c) Run the validation loop for UNTRAINED_NET against the sanity validation dataset.

```
In [125]: sanity_loader = DataLoader(sanity_data, batch_size=1, num_workers=NUM_WORKERS)
train_loader = DataLoader(train_data, batch_size=TRAIN_BATCH_SIZE, num_workers=NUM_WORKERS)
val_loader = DataLoader(val_data, batch_size=1, num_workers=NUM_WORKERS)
```

```
Out[125]: 5.59351152662911e-05
```

```
In [126]:
```

```
Out[126]: <torch.utils.data.DataLoader at 0x15d206f70>
```

```
In [127]:
```

```
Out[127]: <torch.utils.data.DataLoader at 0x15d074bb0>
```

```
In [128]: %%time
```

```
mIoU over the sanity dataset:6.727359952357058e-05
CPU times: user 793 ms, sys: 247 ms, total: 1.04 s
Wall time: 540 ms
```

Problem 5

(a) Define an optimizer to train the given loss function.

Feel free to choose your optimizer of choice from [\(https://pytorch.org/docs/stable/optim.html\)](https://pytorch.org/docs/stable/optim.html).

```
In [129]: from torch import optim
def get_optimizer(net):
    # Write me
    return optim.SGD(net.parameters(), lr=0.006, momentum=0.9)
```

(b) Write the training loop to train the network.

```
In [130]: def train(train_loader, net, optimizer, loss_graph):
    main_loss = 0
    for i, data in enumerate(train_loader):

        inputs, masks = data

        if USE_GPU:
            inputs = inputs.cuda()
            masks = masks.cuda()
            net = net.cuda()

        optimizer.zero_grad()
        main_loss = net(inputs, gts=masks)

        loss_graph.append(main_loss.item())
        main_loss.backward()
        optimizer.step()

    return main_loss
```

(c) Create OVERFIT_NET and train it on the single image dataset.

Single image training is helpful for debugging and hyper-parameter tuning (e.g. learning rate, etc.) as it is fast even on a single CPU. In particular, you can work with a single image until your loss function is consistently decreasing during training loop and the network starts producing a reasonable output for this training image. Training on a single image also teaches about overfitting, particularly when comparing it with more thorough forms of network training.

```
In [116]: from torch import optim
def get_optimizer(net):
    # Write me
```

In [76]:

```
%time
%matplotlib notebook

# The whole training on a single image (20-40 epochs) should take only
# Below we create a (deep) copy of untrained_net and train it on a single image
# Later, we will create a separate (deep) copy of untrained_net to be used for inference
# NOTE: Normally, one can create a new net via declaration new_net = MyNet(). This is fine.
# But if two nets are declared that way creates *different* untrained nets. This notebook shows that.
overfit_net = copy.deepcopy(untrained_net)

# set loss function for the net
overfit_net.criterion = nn.CrossEntropyLoss(ignore_index=255)
#untrained_net.criterion = MyCrossEntropyLoss(ignore_index=255)

# You can change the number of EPOCHS
EPOCH = 40

# switch to train mode (original untrained_net was set to eval mode)
overfit_net.train()

optimizer = get_optimizer(overfit_net)

print("Starting Training...")

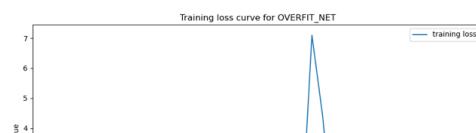
loss_graph = []

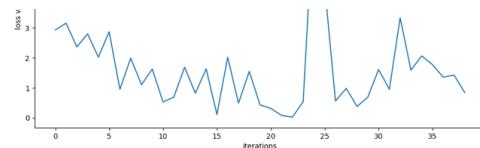
fig = plt.figure(figsize=(12,6))
plt.subplots_adjust(bottom=0.2,right=0.85,top=0.95)
ax = fig.add_subplot(1,1,1)

for e in range(EPOCH):
    loss = train(sanity_loader, overfit_net, optimizer, loss_graph)
    ax.clear()
    ax.set_xlabel('iterations')
    ax.set_ylabel('loss value')
    ax.set_title('Training loss curve for OVERFIT_NET')
    ax.plot(loss_graph, label='training loss')
    ax.legend(loc='upper right')
    fig.canvas.draw()
    print("Epoch: {} Loss: {}".format(e, loss))

%matplotlib inline
```

Starting Training...





```
Epoch: 0 Loss: 2.932879686355591
Epoch: 1 Loss: 3.161799669265747
Epoch: 2 Loss: 2.3720486164093018
Epoch: 3 Loss: 2.8044686317443848
Epoch: 4 Loss: 2.0266449451446533
Epoch: 5 Loss: 2.873180627822876
Epoch: 6 Loss: 0.9573882818222046
Epoch: 7 Loss: 1.9959110021591187
Epoch: 8 Loss: 1.1064273118972778
Epoch: 9 Loss: 1.6312289237976074
Epoch: 10 Loss: 0.5308806300163269
Epoch: 11 Loss: 0.6939668655395508
Epoch: 12 Loss: 1.6909868717193604
Epoch: 13 Loss: 0.8248165249824524
Epoch: 14 Loss: 1.6365500688552856
Epoch: 15 Loss: 0.1126379668712616
Epoch: 16 Loss: 2.024524450302124
Epoch: 17 Loss: 0.49278029799461365
Epoch: 18 Loss: 1.5506278276443481
Epoch: 19 Loss: 0.4351884424686432
Epoch: 20 Loss: 0.31585443019866943
Epoch: 21 Loss: 0.08498293161392212
Epoch: 22 Loss: 0.022573504596948624
Epoch: 23 Loss: 0.5425437688827515
Epoch: 24 Loss: 7.097770690917969
Epoch: 25 Loss: 4.38802433013916
Epoch: 26 Loss: 0.5628655552864075
Epoch: 27 Loss: 0.9860445261001587
Epoch: 28 Loss: 0.38245636224746704
Epoch: 29 Loss: 0.6915167570114136
Epoch: 30 Loss: 1.6142913103103638
Epoch: 31 Loss: 0.9541700482368469
Epoch: 32 Loss: 3.3415770530700684
Epoch: 33 Loss: 1.5963431596755981
Epoch: 34 Loss: 2.067178249359131
Epoch: 35 Loss: 1.7778452634811401
```

```
Epoch: 36 Loss: 1.357385516166687
Epoch: 37 Loss: 1.4308255910873413
Epoch: 38 Loss: 0.8413667678833008
Epoch: 39 Loss: 0.8404616713523865
CPU times: user 1min 27s, sys: 23.8 s, total: 1min 51s
Wall time: 58 s
```

Qualitative and quantitative evaluation of predictions (untrained vs overfit nets) - fully implemented.

In [77]:

```

# switch back to evaluation mode
overfit_net.eval()

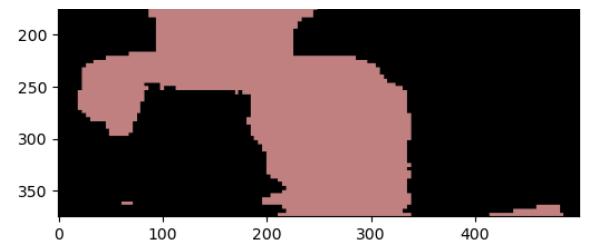
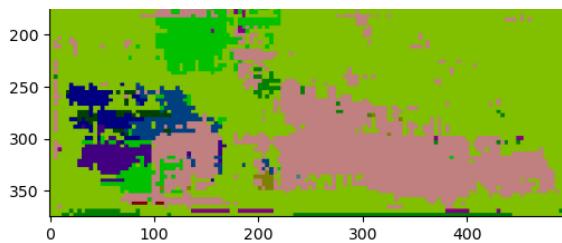
sample_img, sample_target = Normalize(*norm)(*JointToTensor())(sample1
if USE_GPU:
    sample_img = sample_img.cuda()
sample_output_0 = overfit_net.forward(sample_img[None])
sample_output_U = untrained_net.forward(sample_img[None])

# computing mIoU (quantitative measure of accuracy for network predict
if USE_GPU:
    pred_0 = torch.argmax(sample_output_0, dim=1).cpu().numpy()[0]
    pred_U = torch.argmax(sample_output_U, dim=1).cpu().numpy()[0]
else:
    pred_0 = torch.argmax(sample_output_0, dim=1).numpy()[0]
    pred_U = torch.argmax(sample_output_U, dim=1).numpy()[0]

gts = torch.from_numpy(np.array(sample1[1].convert('P'), dtype=np.int3
gts[gts == 255] = -1
conf_0 = eval_semantic_segmentation(pred_0[None], gts[None])
conf_U = eval_semantic_segmentation(pred_U[None], gts[None])

```





```
In [78]: sample_img, sample_target = Normalize(*norm)(*JointToTensor())(sample2
if USE_GPU:
    sample_img = sample_img.cuda()
sample_output_0 = overfit_net.forward(sample_img[None])
sample_output_U = untrained_net.forward(sample_img[None])

# computing mIoU (quantitative measure of accuracy for network predict
if USE_GPU:
    pred_0 = torch.argmax(sample_output_0, dim=1).cpu().numpy()[0]
    pred_U = torch.argmax(sample_output_U, dim=1).cpu().numpy()[0]
else:
    pred_0 = torch.argmax(sample_output_0, dim=1).numpy()[0]
    pred_U = torch.argmax(sample_output_U, dim=1).numpy()[0]

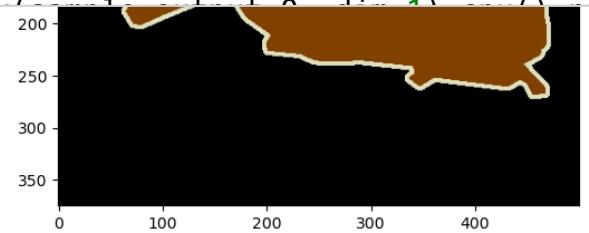
gts = torch.from_numpy(np.array(sample2[1].convert('P'), dtype=np.int3
gts[gts == 255] = -1
conf_0 = eval_semantic_segmentation(pred_0[None], gts[None])
conf_U = eval_semantic_segmentation(pred_U[None], gts[None])

fig = plt.figure(figsize=(14,10))
ax1 = fig.add_subplot(2,2,1)
plt.title('image sample')
ax1.imshow(sample2[0])
ax2 = fig.add_subplot(2,2,2)
plt.title('ground truth (target)')
ax2.imshow(sample2[1])
ax3 = fig.add_subplot(2,2,3)
plt.title('UNTRAINED_NET prediction')
ax3.text(10, 25, 'mIoU = {:.>8.6f}'.format(conf_U['miou']), fontsize=2
ax3.imshow(colorize_mask(torch.argmax(sample_output_U, dim=1).cpu().nu
ax4 = fig.add_subplot(2,2,4)
plt.title('OVERFIT_NET prediction (for image it has not seen)')
ax4.text(10, 25, 'mIoU = {:.>8.6f}'.format(conf_0['miou']), fontsize=2
ax4.imshow(colorize_mask(torch.argmax(sample_output_0, dim=1).cpu().nu
ax4.set_title('OVERFIT_NET prediction (for image it has not seen)')

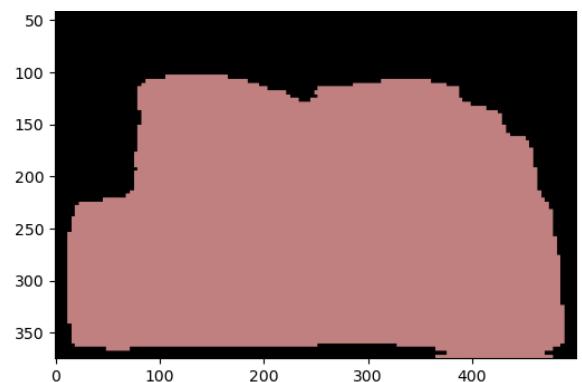
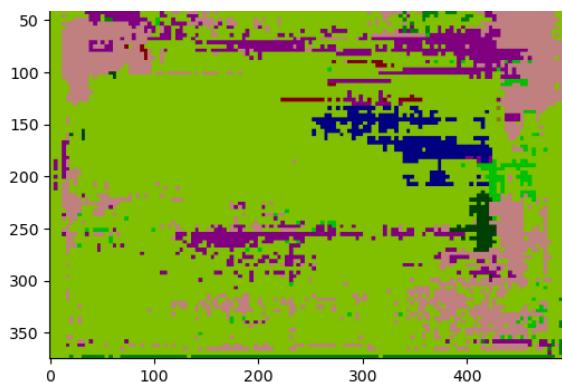
# UNTRAINED_NET prediction
mIoU = 0.000000
# OVERFIT_NET prediction (for image it has not seen)
mIoU = 0.173442
```



UNTRAINED_NET prediction
mIoU = 0.000000



OVERFIT_NET prediction (for image it has not seen)
mIoU = 0.173442



Run the validation loop for OVERFIT_NET against the sanity dataset (an image it was trained on) - fully implemented

In [118]:

```
%time  
mIoU for OVERFIT_NET over its training image:0.7700754988910916  
CPU times: user 832 ms, sys: 238 ms, total: 1.07 s  
Wall time: 624 ms
```

WARNING: For the remaining part of the assignment (below) it is advisable to switch to GPU mode as running each validation and training loop on the whole training set takes over an hour on CPU (there are several such loops below). Note that GPU mode is helpful only if you have a sufficiently good NVIDIA gpu (not older than 2-3 years) and cuda installed on your computer. If you do not have a sufficiently good graphics card available, you can still finish the remaining part in CPU mode (takes a few hours), as the cells below are mostly implemented and test your code written and debugged in the earlier parts above. You can also switch to Google Colaboratory to run the remaining parts below.

You can use validation-data experiments below to tune your hyper-parameters. Normally, validation data is used exactly for this purpose. For actual competitions, testing data is not public and you can not tune hyper-parameters on it.

(d) Evaluate UNTRAINED_NET and OVERFIT_NET on validation dataset.

Run the validation loop for UNTRAINED_NET against the validation dataset:

In [119]:

Out[119]: <torch.utils.data.dataloader.DataLoader at 0x15cf61220>

In [120]:

```
for i, data in enumerate(val_loader):
    print(i)
    print(data)
```

[-0.5826, -0.5826, -0.6001, ..., -0.2150, -0.2325, -0.2325
],
[[0.9319, 0.9319, 0.9319, ..., 0.9668, 0.9842, 0.9842
],
[0.9319, 0.9319, 0.9319, ..., 0.8971, 0.9145, 0.9145
],
[0.9319, 0.9319, 0.9319, ..., 0.7751, 0.7925, 0.7925
],
...,
[-0.6890, -0.6890, -0.7064, ..., -0.0092, -0.0790, -0.0964
],
[-0.6715, -0.6715, -0.6890, ..., 0.1128, 0.0779, 0.0605
],
[-0.6715, -0.6715, -0.6890, ..., 0.1999, 0.1825, 0.1651
]]]), tensor([[[255, 255, 255, ..., 0, 0, 0],
[255, 255, 255, ..., 0, 0, 0],
[255, 255, 255, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 5, 5, 5],

In [*]:

```
%%time
# This will be slow on CPU (around 1 hour or more). On GPU it should take less time.
print("mIoU for UNTRAINED_NET over the entire dataset:{}".format(valida
```

Run the validation loop for OVERFIT_NET against the validation dataset (it has not seen):

In [*]:

```
%%time
# This will be slow on CPU (around 1 hour or more). On GPU it should take less time.
print("mIoU for OVERFIT_NET over the validation dataset:{}".format(val
```

(e) Explain in a few sentences the quantitative results observed in (c) and (d):

Student answer:

in c) The overfit net can only separate areas of the image with high color difference as shown in the old woman. It is not able to separate her hair and face from the background yet able to draw out the region where her sweater occupies in the image. The untrained net is just random garbage.

While in d) the overtrained net confuses the lighter fur of the lama with the higher intensity lighting on the ground.

(f) Create TRAINED_NET and train it on the full training dataset:

In [71]:

```
# Increase TRAIN_BATCH_SIZE if you are using GPU to speed up training.  
# When batch size changes, the learning rate may also need to be adjusted.  
# Note that batch size maybe limited by your GPU memory, so adjust if necessary.  
TRAIN_BATCH_SIZE = 4  
  
# If you are NOT using Windows, set NUM_WORKERS to anything you want,  
# but Windows has issues with multi-process dataloaders, so NUM_WORKERS = 0  
  
sanity_loader = DataLoader(sanity_data, batch_size=1, num_workers=NUM_WORKERS)  
train_loader = DataLoader(train_data, batch_size=TRAIN_BATCH_SIZE, num_workers=NUM_WORKERS)
```

In [72]:

```
%%time
%matplotlib notebook

# This training will be very slow on a CPU (>1hour per epoch). Ideally
# taking only a few minutes per epoch (depending on your GPU and batch
# it is highly advisable that you first finish debugging your net code
# reasonably, e.g. its loss monotonically decreases during training and
# Below we create another (deep) copy of untrained_net. Unlike OVERFIT
trained_net = copy.deepcopy(untrained_net)

# set loss function for the net
trained_net.criterion = nn.CrossEntropyLoss(ignore_index=255)
#trained_net.criterion = MyCrossEntropyLoss(ignore_index=255)

# You can change the number of EPOCHS below. Since each epoch for TRAINED_NET
# the number of required epochs could be smaller compared to OVERFIT_NET
EPOCH = 2

# switch to train mode (original untrained_net was set to eval mode)
trained_net.train()

optimizer = get_optimizer(trained_net)

print("Starting Training...")

loss_graph = []

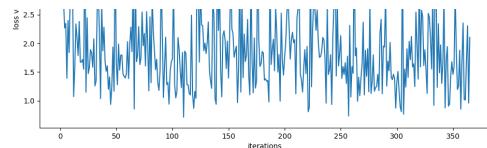
fig = plt.figure(figsize=(12,6))
plt.subplots_adjust(bottom=0.2,right=0.85,top=0.95)
ax = fig.add_subplot(1,1,1)

for e in range(EPOCH):
    loss = train(train_loader, trained_net, optimizer, loss_graph)
    ax.clear()
    ax.set_xlabel('iterations')
    ax.set_ylabel('loss value')
    ax.set_title('Training loss curve for TRAINED_NET')
    ax.plot(loss_graph, label='training loss')
    ax.legend(loc='upper right')
    fig.canvas.draw()
    print("Epoch: {} Loss: {}".format(e, loss))

%matplotlib inline
```

Starting Training...





```
Epoch: 0 Loss: 2.103353977203369
```

```
Epoch: 1 Loss: 2.137852668762207
```

```
CPU times: user 1h 43min 40s, sys: 15min 30s, total: 1h 59min 11s
```

```
Wall time: 1h 11min 49s
```

(g) Qualitative and quantitative evaluation of predictions (OVERFIT_NET vs TRAINED_NET):

In [73]:

```

# switch back to evaluation mode
trained_net.eval()

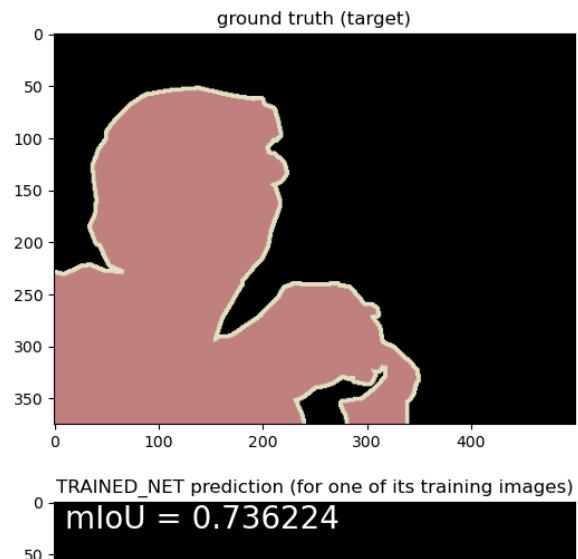
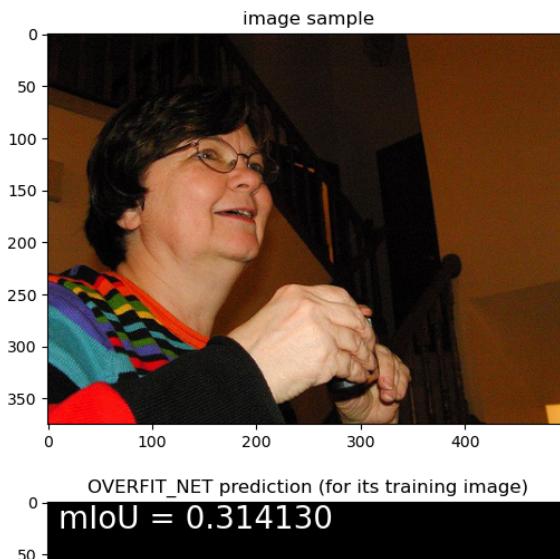
sample_img, sample_target = Normalize(*norm)(*JointToTensor())(*sample1
if USE_GPU:
    sample_img = sample_img.cuda()
sample_output_0 = overfit_net.forward(sample_img[None])
sample_output_T = trained_net.forward(sample_img[None])

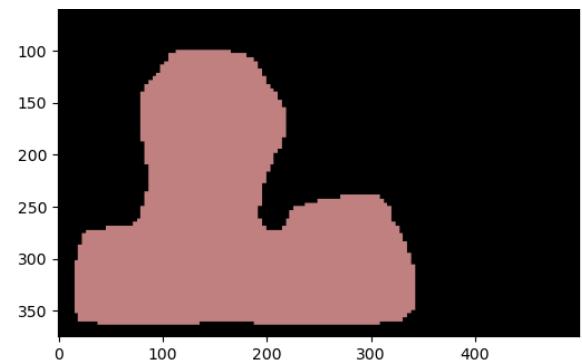
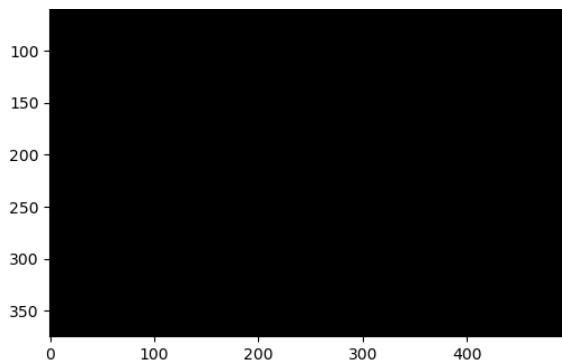
# computing mIoU (quantitative measure of accuracy for network predict
pred_T = torch.argmax(sample_output_T, dim=1).cpu().numpy()[0]
pred_0 = torch.argmax(sample_output_0, dim=1).cpu().numpy()[0]
gts = torch.from_numpy(np.array(sample1[1].convert('P'), dtype=np.int32))
gts[gts == 255] = -1
conf_T = eval_semantic_segmentation(pred_T[None], gts[None])
conf_0 = eval_semantic_segmentation(pred_0[None], gts[None])

fig = plt.figure(figsize=(14,10))
ax1 = fig.add_subplot(2,2,1)
plt.title('image sample')
ax1.imshow(sample1[0])
ax2 = fig.add_subplot(2,2,2)
plt.title('ground truth (target)')
ax2.imshow(sample1[1])
ax3 = fig.add_subplot(2,2,3)
plt.title('OVERFIT_NET prediction (for its training image)')
ax3.text(10, 25, 'mIoU = {:.2f}'.format(conf_0['miou']), fontsize=20)
ax3.imshow(colorize_mask(torch.argmax(sample_output_0, dim=1).cpu().numpy()))
ax4 = fig.add_subplot(2,2,4)
plt.title('TRAINED_NET prediction (for one of its training images)')
ax4.text(10, 25, 'mIoU = {:.2f}'.format(conf_T['miou']), fontsize=20)
ax4.imshow(colorize_mask(torch.argmax(sample_output_T, dim=1).cpu().numpy()))

```

Out[73]: <matplotlib.image.AxesImage at 0x29bfbd550>



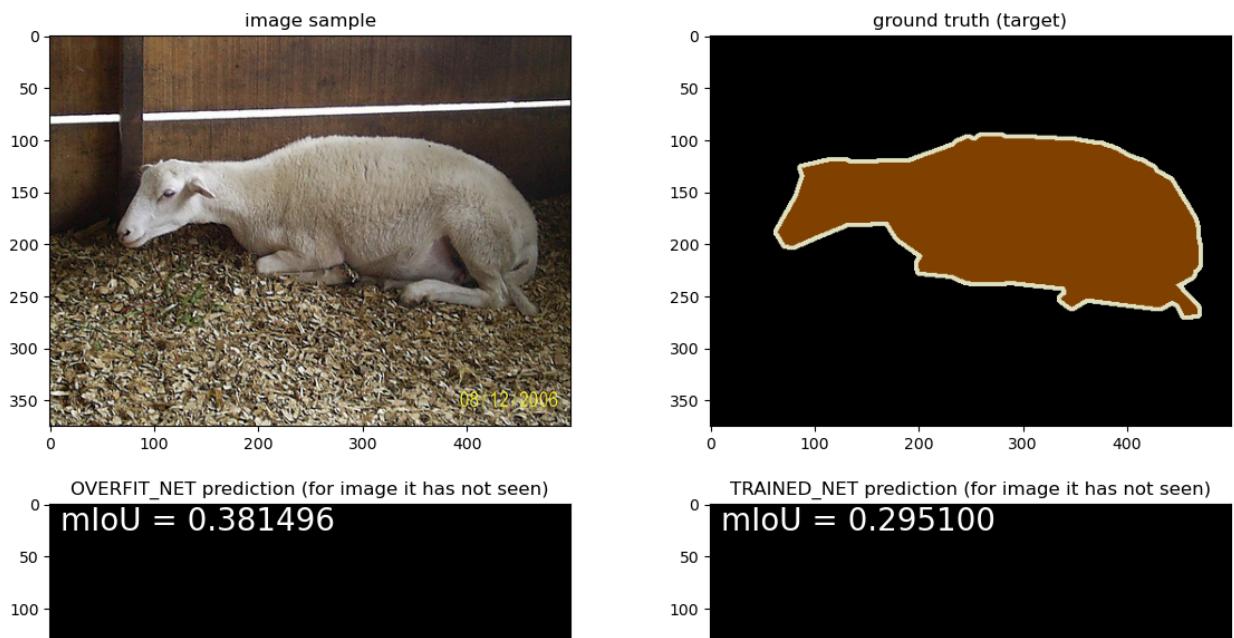


```
In [74]: sample_img, sample_target = Normalize(*norm)(*JointToTensor())(*sample2
if USE_GPU:
    sample_img = sample_img.cuda()
sample_output_0 = overfit_net.forward(sample_img[None])
sample_output_T = trained_net.forward(sample_img[None])

# computing mIoU (quantitative measure of accuracy for network predict
pred_0 = torch.argmax(sample_output_0, dim=1).cpu().numpy()[0]
pred_T = torch.argmax(sample_output_T, dim=1).cpu().numpy()[0]
gts = torch.from_numpy(np.array(sample2[1].convert('P'), dtype=np.int32))
gts[gts == 255] = -1
conf_0 = eval_semantic_segmentation(pred_0[None], gts[None])
conf_T = eval_semantic_segmentation(pred_T[None], gts[None])

fig = plt.figure(figsize=(14,10))
ax1 = fig.add_subplot(2,2,1)
plt.title('image sample')
ax1.imshow(sample2[0])
ax2 = fig.add_subplot(2,2,2)
plt.title('ground truth (target)')
ax2.imshow(sample2[1])
ax3 = fig.add_subplot(2,2,3)
plt.title('OVERFIT_NET prediction (for image it has not seen)')
ax3.text(10, 25, 'mIoU = {:.2f}'.format(conf_0['miou']), fontsize=20)
ax3.imshow(colorize_mask(torch.argmax(sample_output_0, dim=1).cpu().numpy()))
ax4 = fig.add_subplot(2,2,4)
plt.title('TRAINED_NET prediction (for image it has not seen)')
ax4.text(10, 25, 'mIoU = {:.2f}'.format(conf_T['miou']), fontsize=20)
ax4.imshow(colorize_mask(torch.argmax(sample_output_T, dim=1).cpu().numpy()))
```

Out[74]: <matplotlib.image.AxesImage at 0x29c138490>



(h) Evaluate TRAINED_NET on validation dataset.

Run the validation loop for TRAINED_NET against the validation dataset (it has not seen):

In [*]:

```
%%time  
# This will be slow on CPU (around 1 hour). On GPU it should take only
```

Problem 6

For the network that you implemented, write a paragraph or two about limitations / bottlenecks about the work. What could be improved? What seems to be some obvious issues with the existing works?

The network I have trained required the images of one consistant size to be used as the input. It is hard to have the neural net to fit the finner details of the sheep. Unfortunately I do not have time to run the verification for the whole training and validation datasets so I can say nothing about the performance limitations of the network.

In []:

In []:

In []:

In []: